

Homework 5

CS211 - Fall 2024

(You should try to answer first and then compare your solution with the ChatGPT solution)

This homework is designed to help you become familiar with scheduling concepts and algorithms in operating systems

1. Let P_1, P_2, P_3, P_4 be processes with execution times of 53, 8, 68, and 24, respectively. Find the average waiting time and average completion time for the following scheduling scenarios: the best-case FCFS, the worst-case FCFS, and Round-Robin with quantum values of 1, 5, 8, 10, and 20.
2. Why does Shortest Remaining Time First (SRTF) scheduling help achieve the best response time and I/O throughput? Provide an example to illustrate your explanation.
3. What is the CPU scheduling policy of Linux? Briefly describe this policy and explain its strengths and weaknesses (e.g., in terms of response time, fairness, overhead, etc.)

① Best case FCFS: P_2, P_4, P_1, P_3

- avg waiting time: $\frac{0 + 8 + (8+24) + (8+24+53)}{4} = 31.25$

- avg completion time: $\frac{8 + (8+24) + (8+24+53) + (8+24+53+68)}{4} = 69.5$

Worst case FCFS : P_3, P_1, P_4, P_2

- avg waiting time: $\frac{0 + 68 + (68+53) + (68+53+24)}{4} = 83.5$

- avg completion time: $\frac{68 + (68+53) + (68+53+24) + (68+53+24+8)}{4} = 121.75$

Round-robin assuming FCFS order is : P_1, P_2, P_3, P_4

RR quantum = 1:

- avg waiting time: $\frac{0 + 1 + 2 + 3}{4} = 1.5$

- avg completion time:
$$\begin{aligned} & (53 + 8 + 52 + 24) \\ & + (8 + 8 + 7 + 7) \\ & + (68 + 53 + 8 + 24) \\ & + (24 + 8 + 24 + 24) \end{aligned} \Bigg| 4 = 100$$

RR quantum = 5:

- avg waiting time: $\frac{0 + 5 + (5+5) + (5+5+5)}{4} = 7.5$

- avg completion time:
$$\begin{aligned} & (53 + 8 + 50 + 24) \\ & + (8 + 10 + 5 + 5) \\ & + (68 + 53 + 8 + 24) \\ & + (24 + 25 + 8 + 25) \end{aligned} \Bigg| 4 = 99.5$$

RR quantum = 8

- avg waiting time: $\frac{0+8+16+24}{4} = 12$

- avg completion time: $(53+8+48+24) + (8+8) + (68+53+8+24) + (24+24+8+24)$

$$4 = 95.5$$

RR quantum = 10:

- avg waiting time: $\frac{0+10+(10+8)+(10+8+10)}{4} = 14$

- avg completion time: $(53+8+50+24) + (10+8) + (68+53+8+24) + (24+30+8+30)$

$$4 \approx 99.5$$

RR quantum = 20:

- avg waiting time: $\frac{0+(20+8)+(20+8+20)}{4} = 19$

- avg completion time: $(53+8+40+24) + (20+8) + (68+53+8+24) + (24+8+40+40)$

$$4 = 104.5$$

② Avg response time:

$$\frac{\sum_i^n (\text{Time job } i \text{ gets started} - \text{Time job } i \text{ enters queue})}{n}$$

Avg I/O throughput:

$$\frac{\text{Number of jobs done}}{\text{Unit of time}}$$

$$\text{or } \frac{\text{Number of all jobs (n)}}{\text{Total time to complete all jobs}}$$

SRTF prioritize shorter jobs → Reduce overall waiting time across tasks on queue

Reduce response time theo cap so công
(only big jobs suffer delay, but many short jobs can reduce waiting for big jobs to get done)

↪ Benefits from this if there are fewer big jobs

Increase the count of jobs done ⇒ Better throughput

Example: P_1, P_2, P_3, P_4 with duration 53, 8, 68, 24

Assume all jobs enter queue at time $t=0$

Avg response time

Avg throughput

Vanilla FCFS

$$\frac{0+53+8+68}{4} = 32.25$$

$$\frac{4}{(0+53)+(53+8)+(53+8+68)+(53+8+68+24)} = \frac{4}{396} \approx 1.0101\%$$

SRTF

$$\frac{0+8+24+53}{4} = 21.25$$

\downarrow
Better than FCFS,
not as good as RR

$$\frac{4}{(0+8)+(8+24)+(8+24+53)+(8+24+53+68)} = \frac{4}{278} \approx 1.4388\% \rightarrow \text{Better than FCFS and RR}$$

Round robin
quantum = 10

$$\frac{0+10+18+28}{4} = 14$$

$$\frac{4}{99.5 \times 4} = \frac{4}{398} \approx 1.005\%$$

\Rightarrow SRTF is good but not outstanding :)

③ Classical alg. with no direct equivalent in linux: SRTF

Classical scheduler	FCFS	Multilevel Queue	RR	Batch scheduling
Linux scheduler	SCHED-FIFO	SCHED-CFS	SCHED-RR	SCHED-BATCH
Response time	Depends (long if big jobs arrive first, short otherwise)	Short (newer jobs has exec. time = 0 \rightarrow higher priority \rightarrow get started soon)	Short (wait for at most quantum duration * queue length)	Long (not focus on optimize response time)
Fairness	Doesn't care	Very fair (ensure fair share of CPU, prioritize jobs with fewer exec. time)	Very fair (ensure fair share of CPU time)	Doesn't care
Overhead	Very little (only a basic queue, no extra context switching)	A lot (Needs a RBTree & priority queue & frequent context switching)	Quite a lot (needs a queue and frequent context switching)	Very little (pre-compute execution order - one off cost)
Throughput	Relatively high (jobs are run to completion, no splitting)	Very low (job gets split, takes forever to done)	Very low (job gets split, takes forever to done)	High (designed to optimize throughput)
Avg - completion time	Short (jobs are run to completion, no split)	Long (jobs get split \rightarrow takes forever)	Long (jobs get split \rightarrow takes forever)	Short (jobs are run to completion, no split)
Avg - waiting time (exec time - queue time)	Long (must wait for previous job to done to start)	Short (jobs get started soon after enter queue)	Short (jobs get started soon after enter queue)	Long (must wait for previous job to done to start)

Linux provides ~~4~~⁶ scheduling algorithms:

- SCHED_FIFO: First come first serve, jobs arrive earlier gets started first, jobs are run to completion
- SCHED_FAIR: Completely fair scheduler, Maintains a priority queue (RBTree), jobs with lesser CPU time gets bump to top of queue → Always select job that's been run the least
- SCHED_RR: Round robin algorithm, CPU time is divided into non-overlapping time slice (quantum), then jobs are rotated to run until done (yield) or timeout (preempt by scheduler)
- SCHED_BATCH: Batch scheduling is often used when we know the list of jobs in advance & can organize them into batches. Once sorted out, each batch is run with little to no scheduling modification.
Jobs are often not split but run to completion.
↳ Suitable for jobs like training models on server.

Above information are extracted from this man page:
man7.org/linux/man-pages/man7/sched.7.html

SCHED_OTHER and SCHED_IDLE are not mentioned because I don't quite understand them (& their use case) at the point of writing this

To find quantum duration of RR scheduling

run this ↗

```
/home/lanphgphm>>> ls /proc/sys/kernel
acct
acpi_video_flags
arch
auto_msgmni
bootloader_type
bootloader_version
bpf_stats_enabled
cad_pid
cap_last_cap
core_pattern
core_pipe_limit
core_uses_pid
ctrl-alt-del
dmesg_restrict
domainname
ftrace_dump_on_oops
ftrace_enabled
hardlockup_all_cpu_backtrace
hardlockup_panic
hostname
hung_task_all_cpu_backtrace
hung_task_check_count
hung_task_check_interval_secs
hung_task_panic
hung_task_timeout_secs
hung_task_warnings
io_delay_type
ioURING_disabled
ioURING_group
kexec_load_disabled
kexec_load_limit_panic
kexec_load_limit_reboot
keys
kptr_restrict
max_lock_depth
max_rcu_stall_to_panic
modprobe
modules_disabled
msg_next_id
msgmax
msgmnb
msgmni
ngroups_max
nmi_watchdog
ns_last_pid
numa_balancing
numa_balancing_promote_rate_limit_MBps
oops_all_cpu_backtrace
oops_limit
osrelease
ostype
overflowgid
overflowuid
panic
panic_on_io_nmi
panic_on_oops
panic_on_rcu_stall
panic_on_unrecoverable_nmi
panic_on_warn
panic_print
perf_cpu_time_max_percent
perf_event_max_contexts_per_stack
perf_event_max_sample_rate
perf_event_max_stack
perf_event_mlock_kb
perf_event_paranoid
pid_max
poweroff_cmd
print-fatal-signals
printk
printk_delay
printk_devkmsg
printk_ratelimit
printk_ratelimit_burst
pty
random
randomize_va_space
real-root-dev
sched_autogroup_enabled
sched_cfs_bandwidth_slice_us
sched_child_runs_first
sched_deadline_period_max_us
sched_deadline_period_min_us
sched_energy_aware
sched_rr_timeslice_ms
sched_rt_period_us
sched_rt_runtime_us
sched_schedstats
sched_util_clamp_max
sched_util_clamp_min
sched_util_clamp_min_rt_default
seccomp
sem
sem_next_id
shm_next_id
shm_rmid_forced
shmll
shmmx
shmmni
soft_watchdog
softlockup_all_cpu_backtrace
softlockup_panic
split_lock_mitigate
stack_tracer_enabled
sysctl_writes_strict
sysrq
tainted
task_delayacct
threads-max
timer_migration
traceoff_on_warning
tracepoint_printk
unknown_nmi_panic
unprivileged_bpf_disabled
unprivileged_userns_clone
user_events_max
usermodehelper
version
warn_limit
watchdog
watchdog_cpumask
watchdog_thresh
yama
```

RR
quantum

Scheduling
specs are here

```
/home/lanphgphm>>> cat /proc/sys/kernel/sched_autogroup_enabled
1
/home/lanphgphm>>> cat /proc/sys/kernel/sched_cfs_bandwidth_slice_us
5000
/home/lanphgphm>>> cat /proc/sys/kernel/sched_child_runs_first
0
/home/lanphgphm>>> cat /proc/sys/kernel/sched_deadline_period_max_us
4194304
/home/lanphgphm>>> cat /proc/sys/kernel/sched_deadline_period_min_us
100
/home/lanphgphm>>> cat /proc/sys/kernel/sched_rr_timeslice_ms
100
/home/lanphgphm>>> cat /proc/sys/kernel/sched_rt_period_us
1000000
/home/lanphgphm>>> cat /proc/sys/kernel/sched_rt_runtime_us
950000
/home/lanphgphm>>> cat /proc/sys/kernel/sched_schedstats
0
/home/lanphgphm>>>
```

quantum
of RR is
100 ms